Tetherline — Visual Polish Review

12 skill scenes × {desktop · tablet · mobile}. Cite issues as **page · scene · viewport** (e.g. "p9 · pipeline · mobile").

pnpm verify: ALL GREEN · typecheck 4/4 · lint 0 fails (130 any-warns) · unit 339/339 · scenes 36/36.
Caveats: grill-screen Playwright capture-timing artifact (p5); mobile chrome dense but functional at 390px.

cite as: **p2 · project-map · <viewport>**

Project map — the home state

WHO/WHY: not a skill and not user-invoked — this is what you land on the moment a repo opens, before any request. The radial map of modules around the project. Each node carries its live Seen/Quiz bars and the header knowledge strip, all rendered from session state. Everything else in this deck is something layered on top of this.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

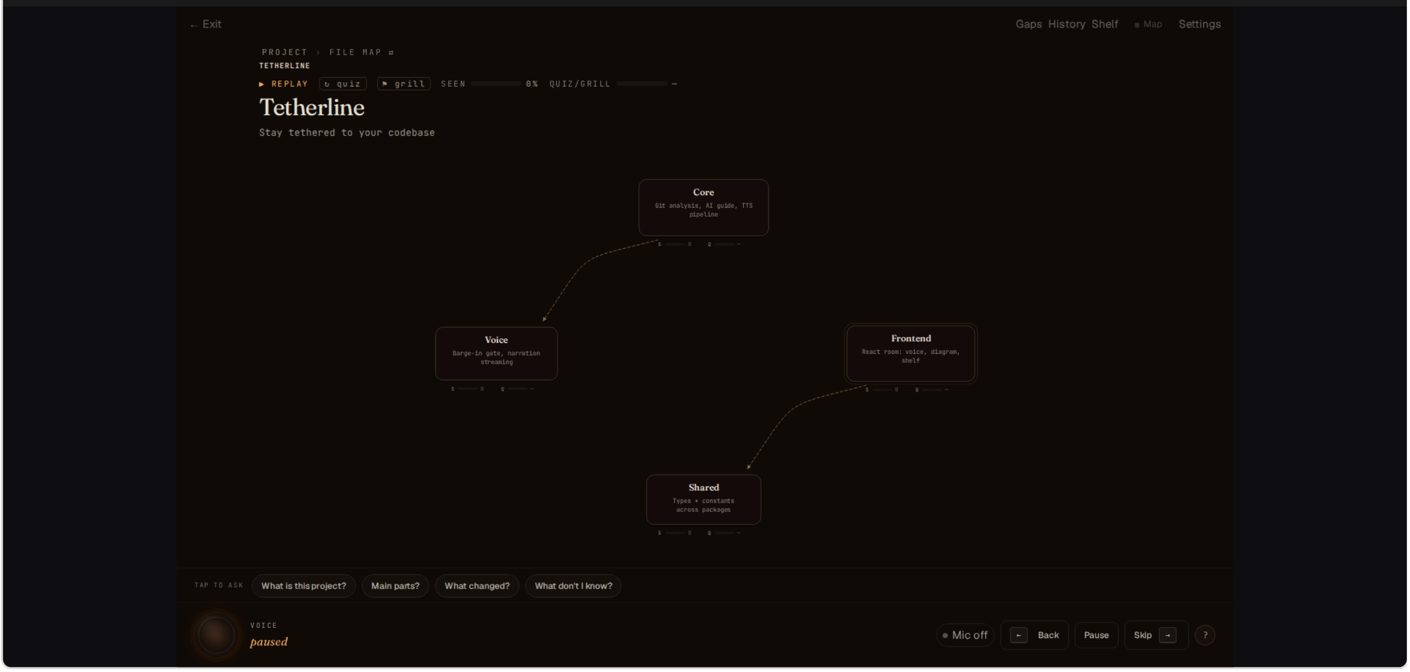

TABLET · 768

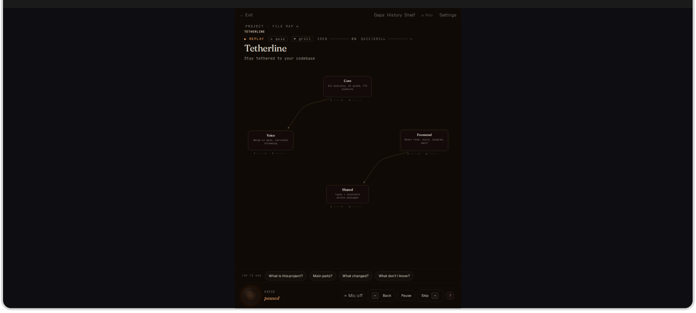

MOBILE · 390

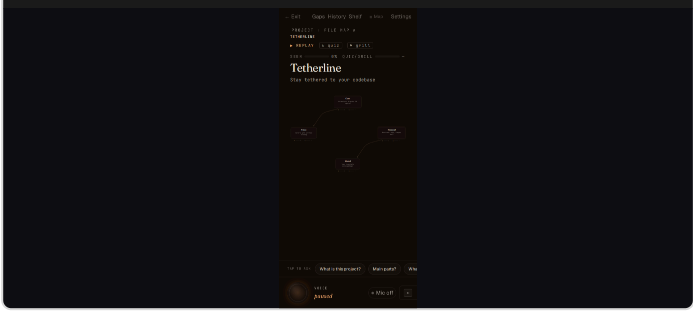

cite as: **p3** · **heatmap** · **<viewport>**

whats_changed — "what you're behind on" (closeable loop)

WHO/WHY: you ask in plain language ("what changed?", "catch me up") — the intent classifier maps that to the whats_changed skill; you never name or click it. Project scope only (drilled-in/per-module drift is a documented follow-up). WHAT: a spoken recap while the diagram washes warm over DRIFT — changed code you haven't caught up on: red (changed, never reviewed) = hot, yellow (changed since you last reviewed) = warm, green (you reviewed it SINCE it changed) = cold, unchanged = no wash. Deliberately NOT raw churn. PATHWAY (now closed): the hot nodes are the agenda — "walk me through the changes" steps each changed area; being walked through one marks its files reviewed-since-change (decision: walked-through = caught up, no quiz gate) and re-emits the heatmap so that node cools to green live. All-green = caught up for the week (the tether re-established). No card; the diagram is the whole visual.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

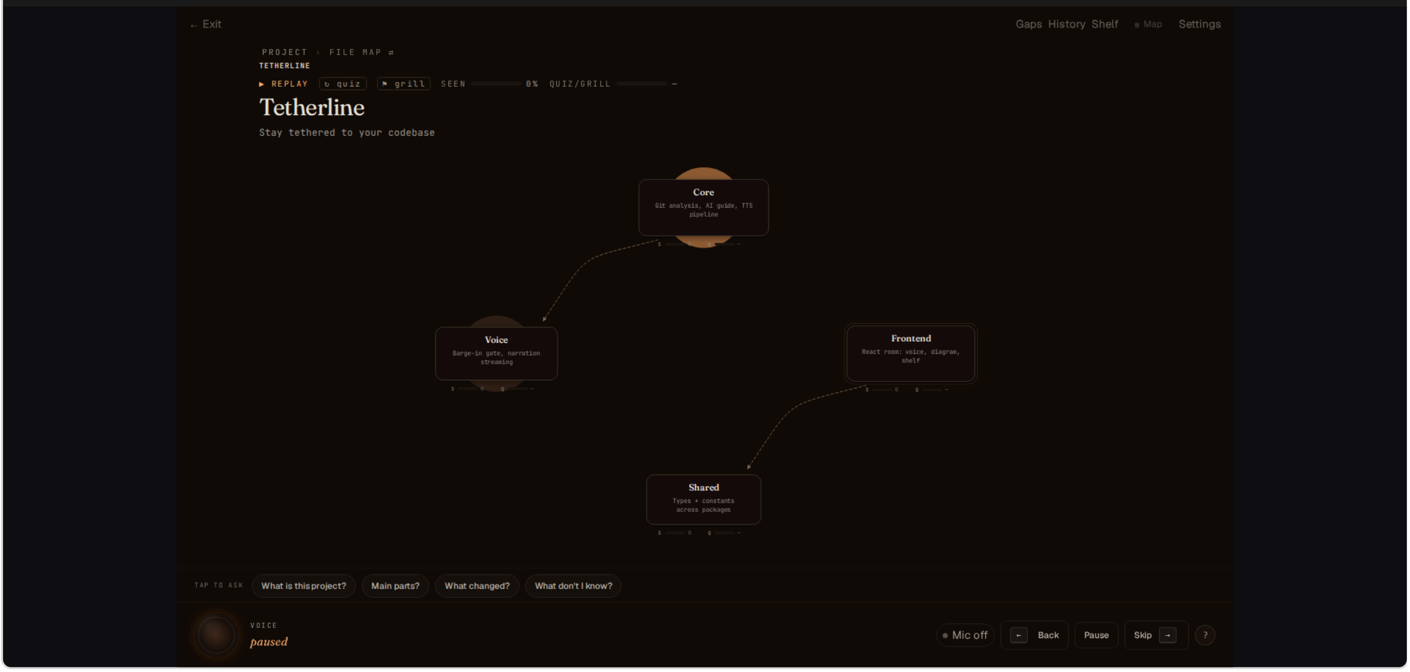

TABLET · 768

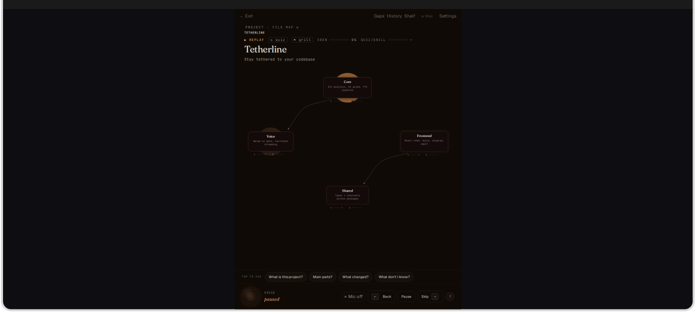

MOBILE · 390

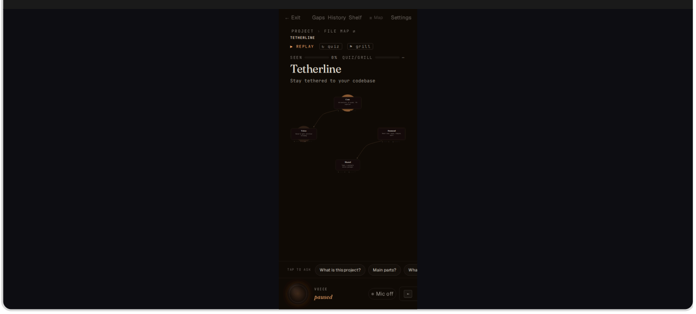

cite as: **p4** · **concern-tint** · **<viewport>**

critique — concern tint

WHO/WHY: you ask for an opinion ("is this good?", "any issues with X?", "what do you think of Y?") — classifier → the critique skill. WHAT: the spoken critique names risky nodes; exactly those glow worry-red. Additive tint, layout unchanged.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

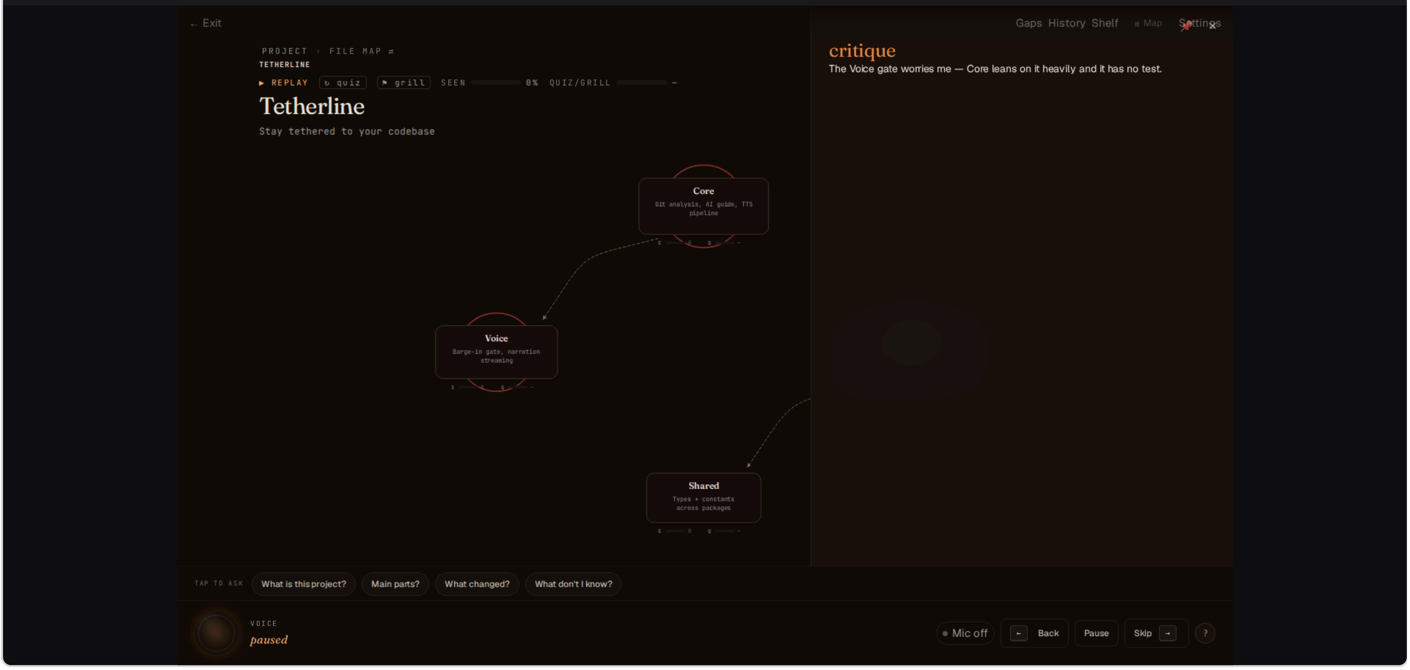

TABLET · 768

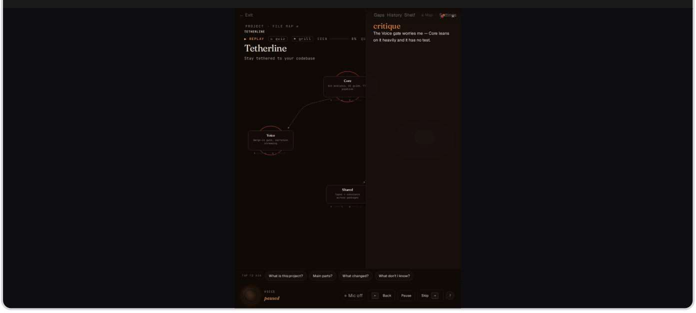

MOBILE · 390

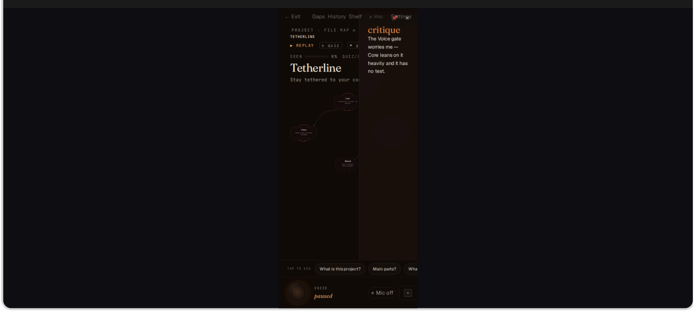

cite as: **p5 · `grill-screen` · `<viewport>`**

grill_me — quiz screen

WHO/WHY: you ask to be tested ("quiz me", "grill me on X", "test my understanding") — classifier → the grill_me skill. WHAT: a calm full-bleed "?" screen replaces the diagram for a Socratic quiz. CAVEAT: Playwright fresh-context capture timing can briefly leak the "grill_me" caption into the screenshot; the live settled DOM is clean (diagram fully replaced). Documented timing artifact, not a live defect.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

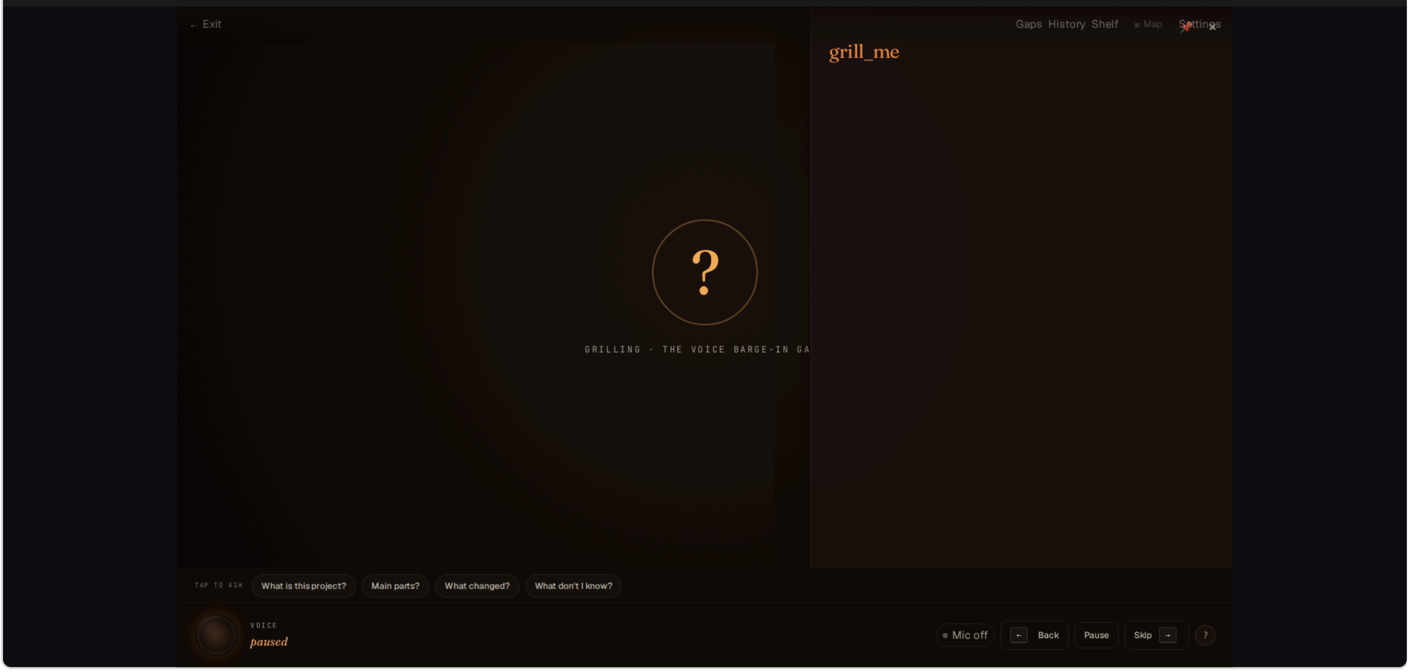

TABLET · 768

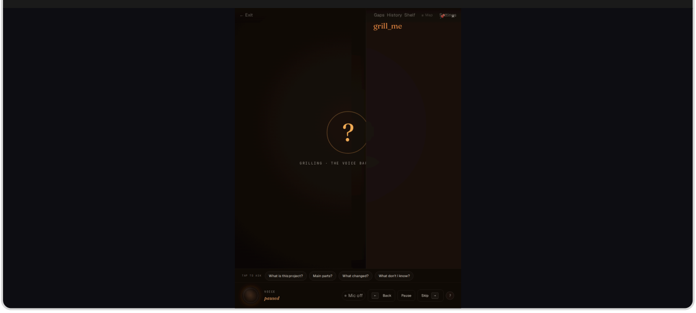

MOBILE · 390

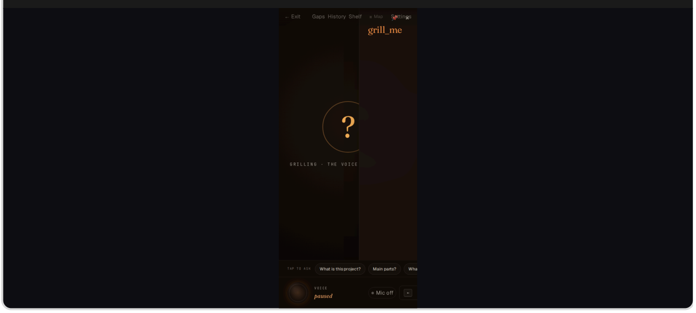

cite as: **p6 · shelf-notes · <viewport>**

Review Shelf · Notebook (annotate)

WHO/WHY: YOU open the Review Shelf from the top chrome; the Notebook tab fills over time as the annotate skill runs ("flag this", "remember this", "note that"). WHAT: saved annotations, each attributed to a module/file. The diagram stays intact behind the shelf.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

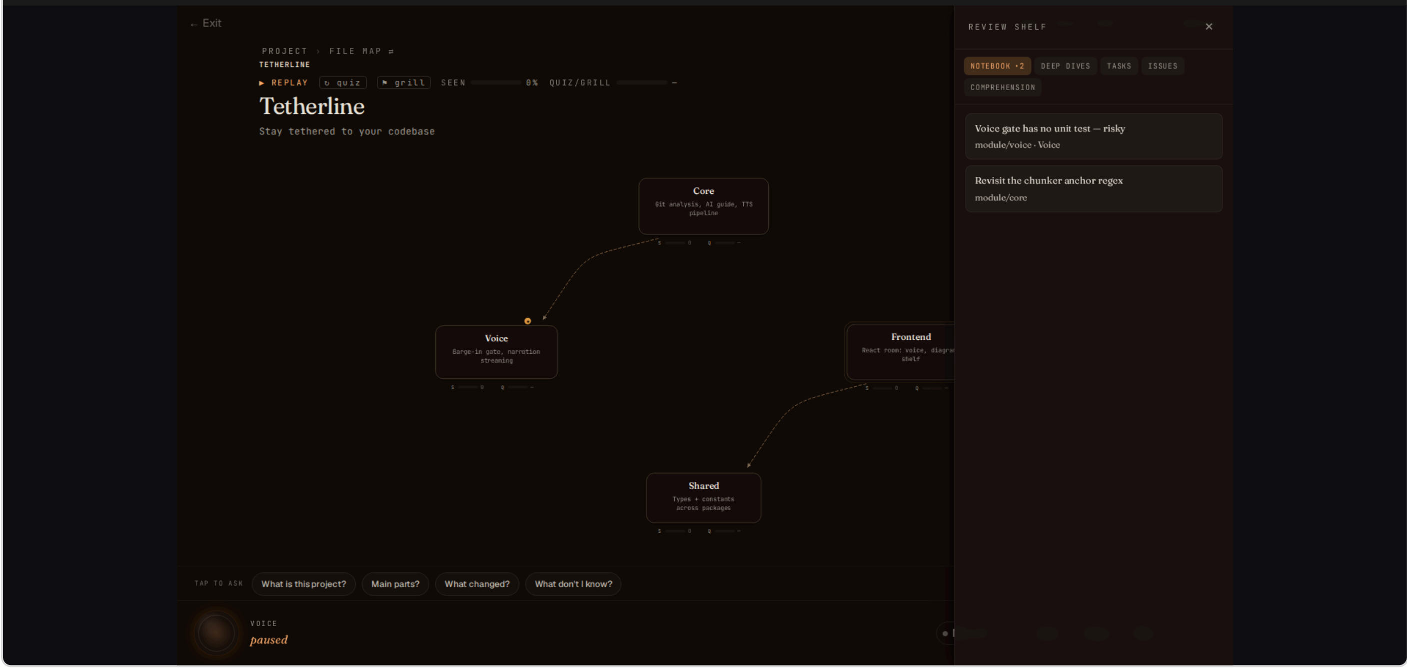

TABLET · 768

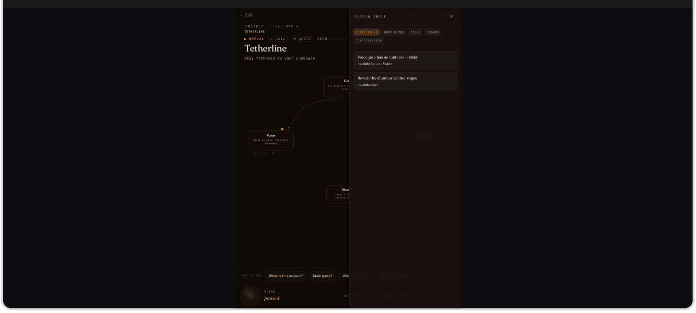

MOBILE · 390

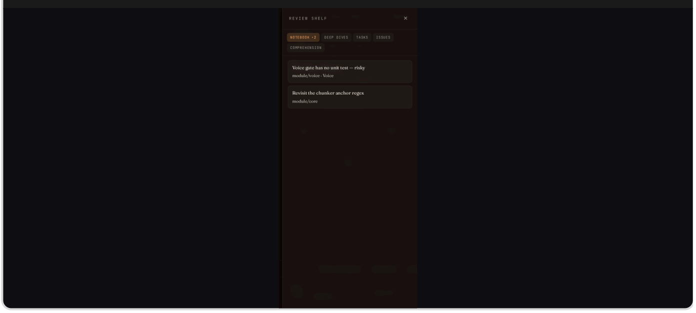

cite as: **p7 · shelf-tasks · <viewport>**

Review Shelf · Tasks (async agent)

WHO/WHY: same shelf, Tasks tab — rows are background agent tasks YOU dispatched ("go fix X", "audit Y"), surfaced here for review. WHAT: task rows with state styling — done · branch:<x> · blocked.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

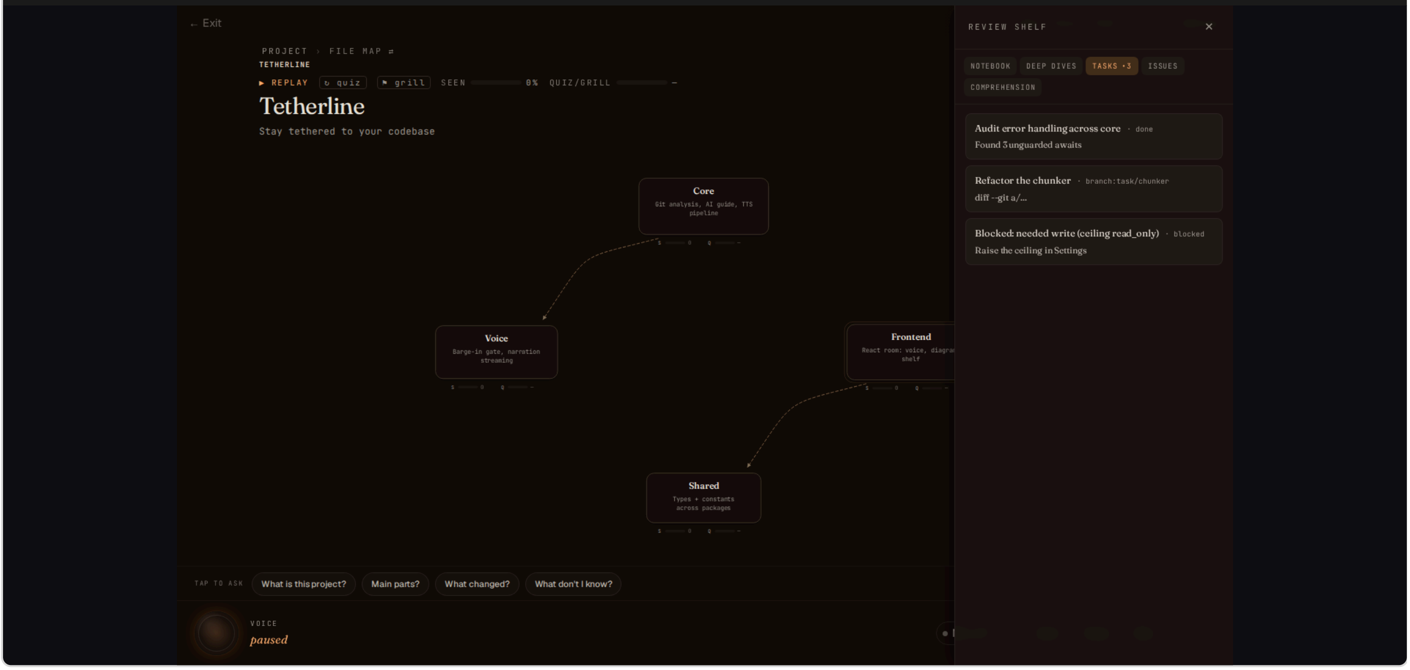

TABLET · 768

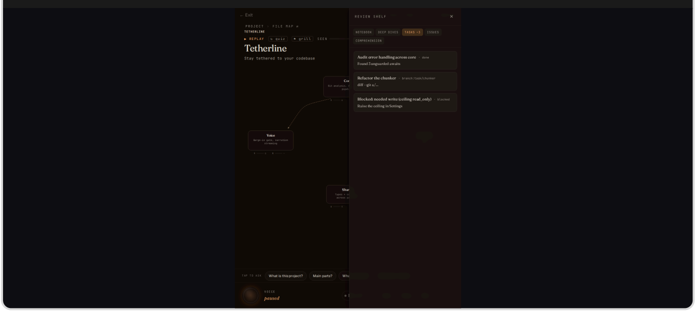

MOBILE · 390

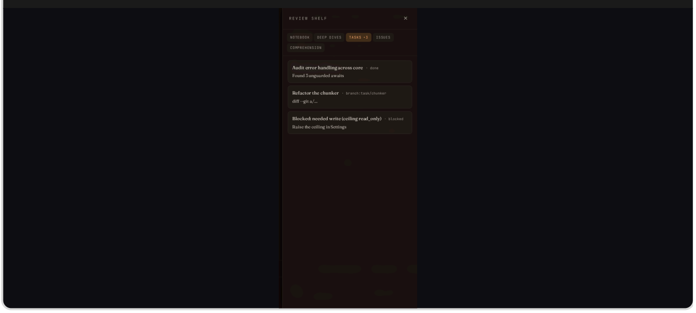

cite as: **p8 · descend · <viewport>**

Drill-in (navigation, not a skill)

WHO/WHY: navigation — you say "show me Core" / "go into X" or click a module node; the view scope-swaps. No skill, no classifier. WHAT: the scoped sub-graph for that module (its files, intra-module edges, retitled header, breadcrumb).

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

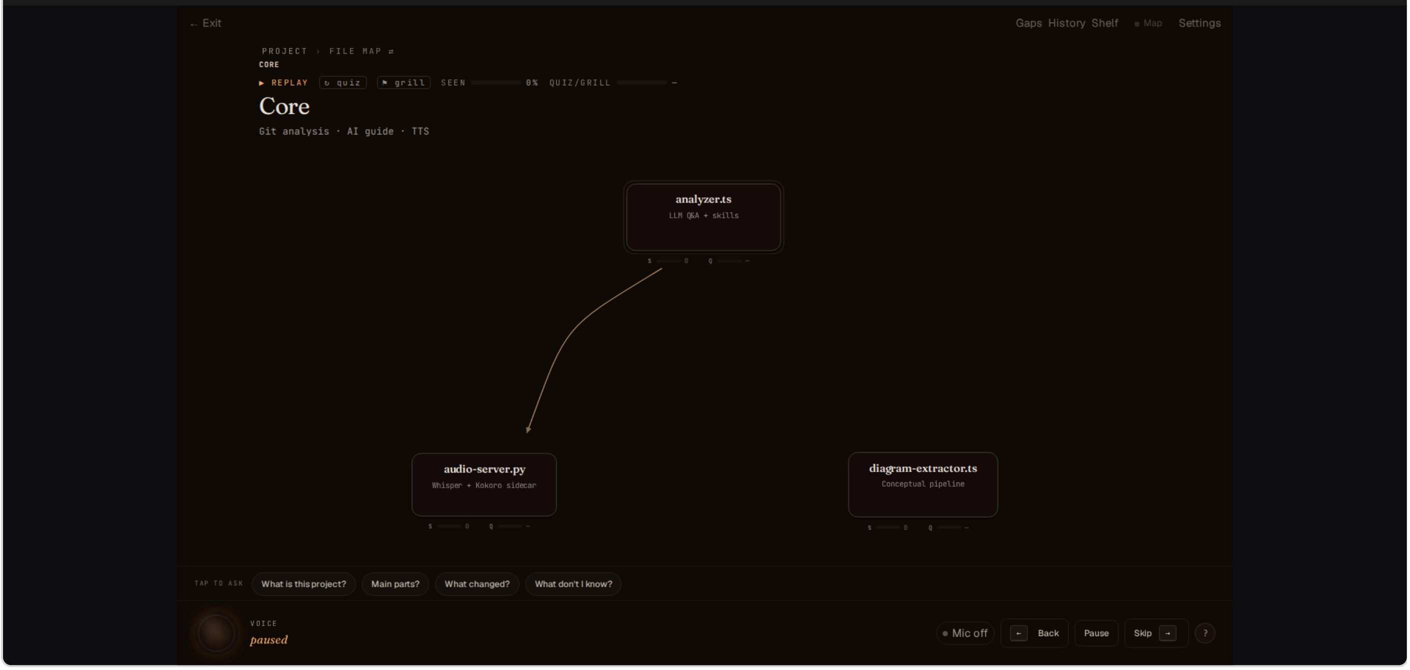

TABLET · 768

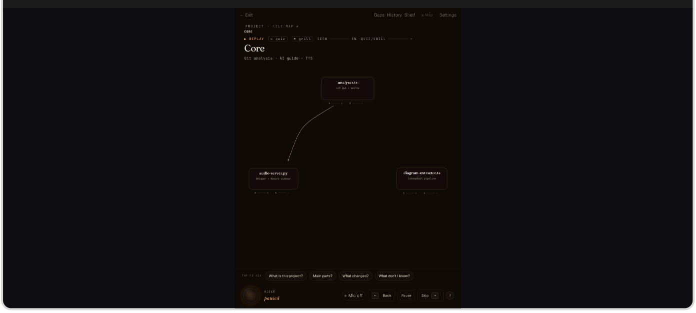

MOBILE · 390

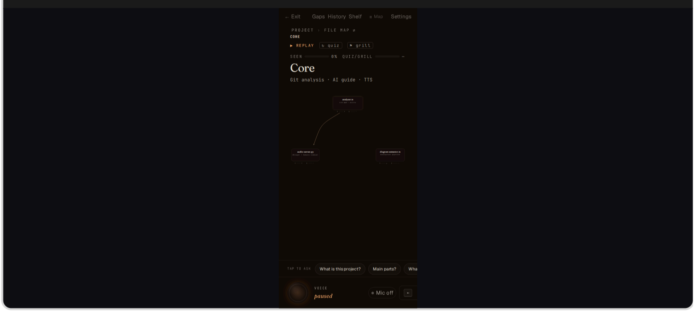

cite as: **p9 · deep-dive · <viewport>**

deep_dive — pocket presentation

WHO/WHY: you ask to go deep ("deep dive on X", "go deeper", "really walk me through this") — classifier → the deep_dive skill. WHAT: a full-bleed ≤10-slide pocket REPLACES the diagram (a sandbox; exiting restores the exact prior canvas). Kicker·focus, N/M cursor, slide rail, exit hint.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

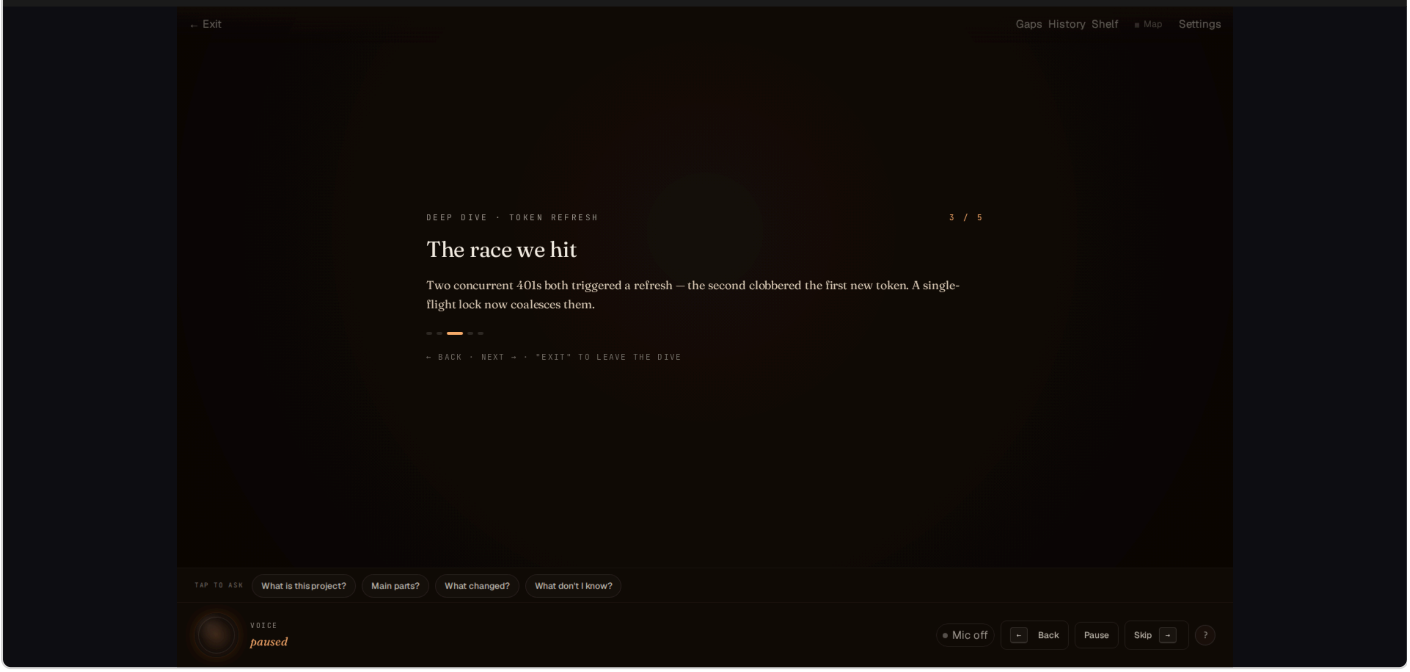

TABLET · 768

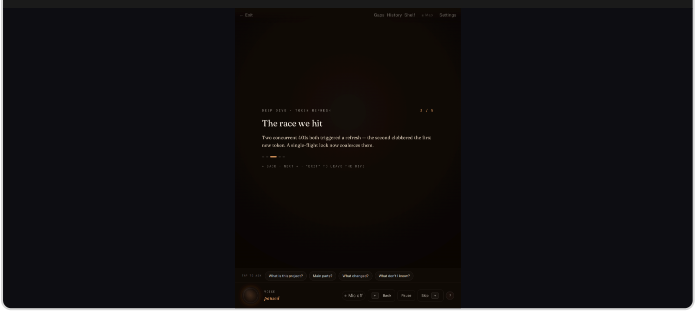

MOBILE · 390

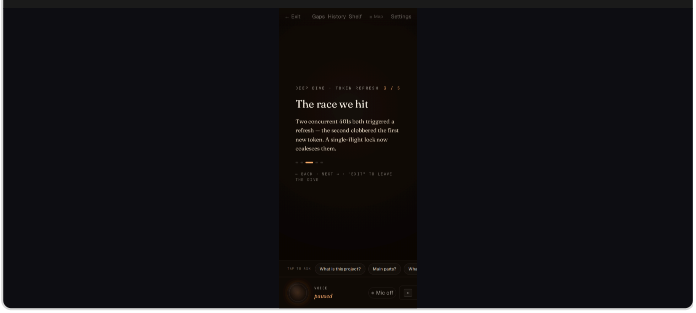

cite as: **p10 · pipeline · <viewport>**

pipeline — data-flow walkthrough

WHO/WHY: you ask how data moves ("show me the data flow", "walk me through how X works end to end") — classifier → the pipeline walkthrough. WHAT: the graph lights one stage at a time source → transform → guard → sink; revealed lit, rest dimmed; header stage strip.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

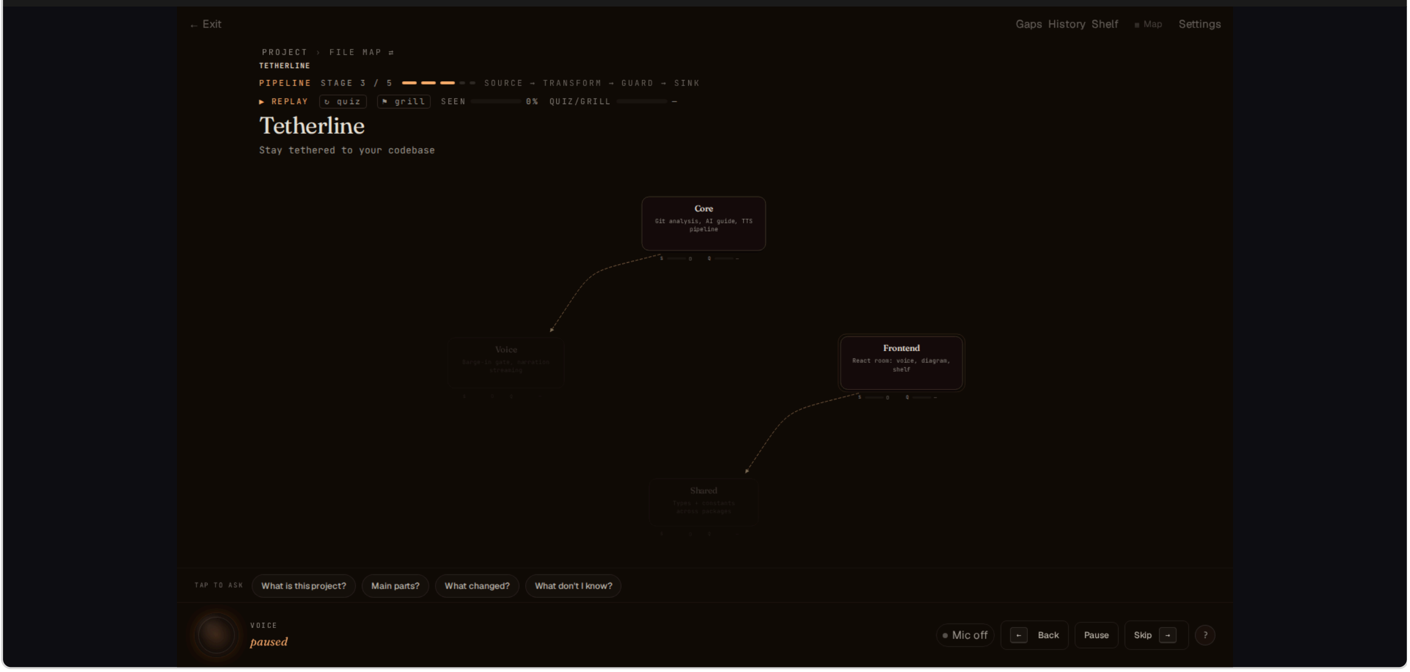

TABLET · 768

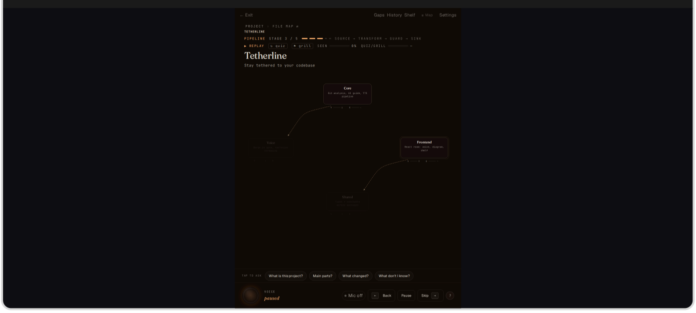

MOBILE · 390

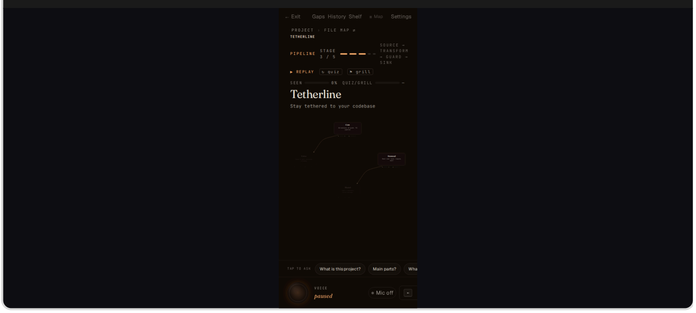

cite as: **p11 · blast-radius · <viewport>**

blast-radius — impact ripple

WHO/WHY: you ask about impact ("what touches X?", "blast radius of Y", "what depends on this?") — classifier → the blast-radius skill. WHAT: a BFS over the import graph from the changed node; concentric rings fade outward by hop distance (hop 0 = solid warm epicenter).

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

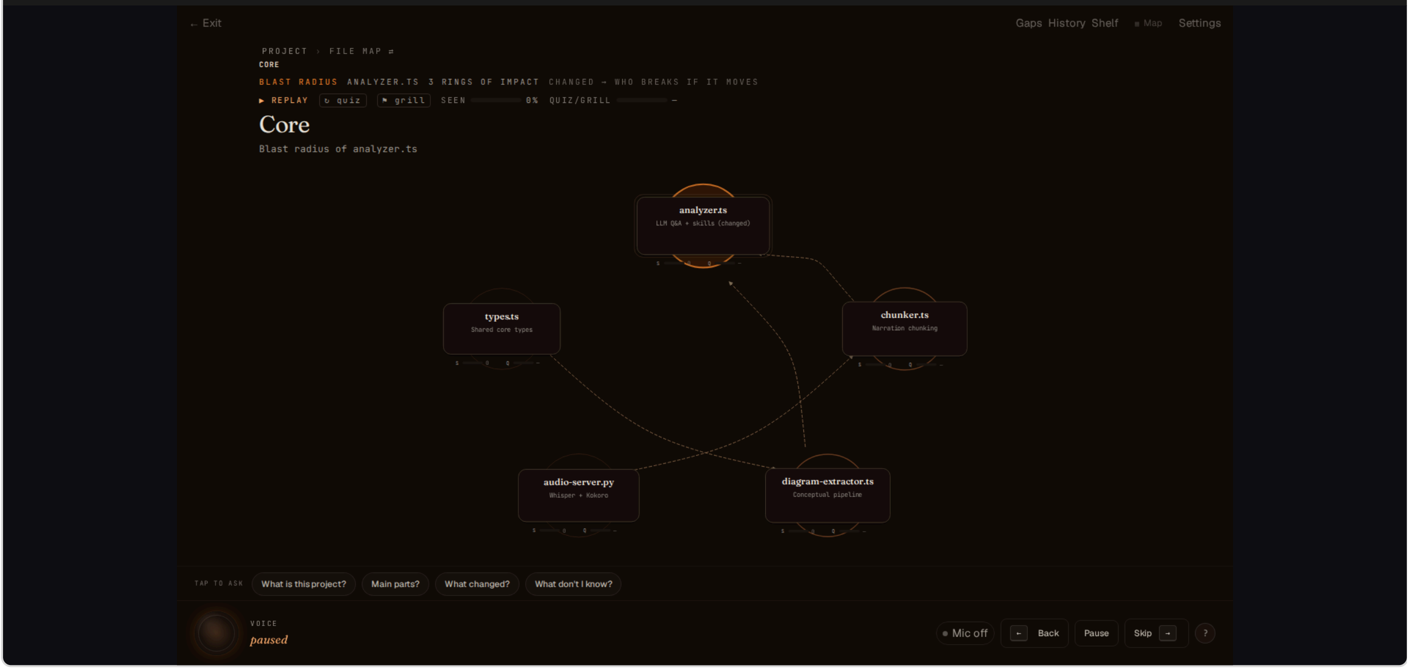

TABLET · 768

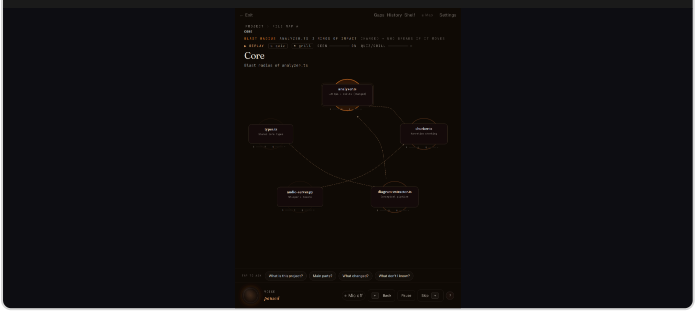

MOBILE · 390

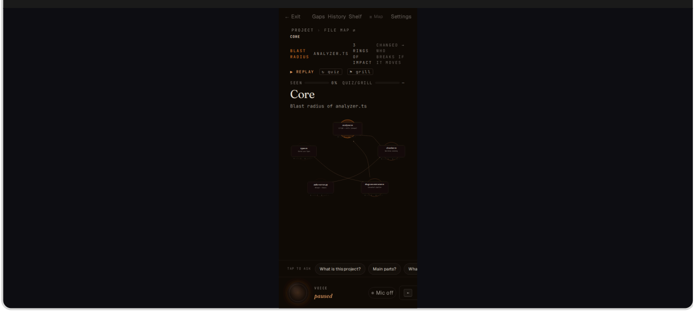

cite as: **p12 · guided-mode · <viewport>**

guided tour — planned walkthrough

WHO/WHY: you ask for a guided tour ("walk me through the project", "give me the tour", "teach me this codebase") — a top-down planned tour drives the session. WHAT: a header spine of covered ▸ current (wide amber) ▸ upcoming steps with the current step name.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

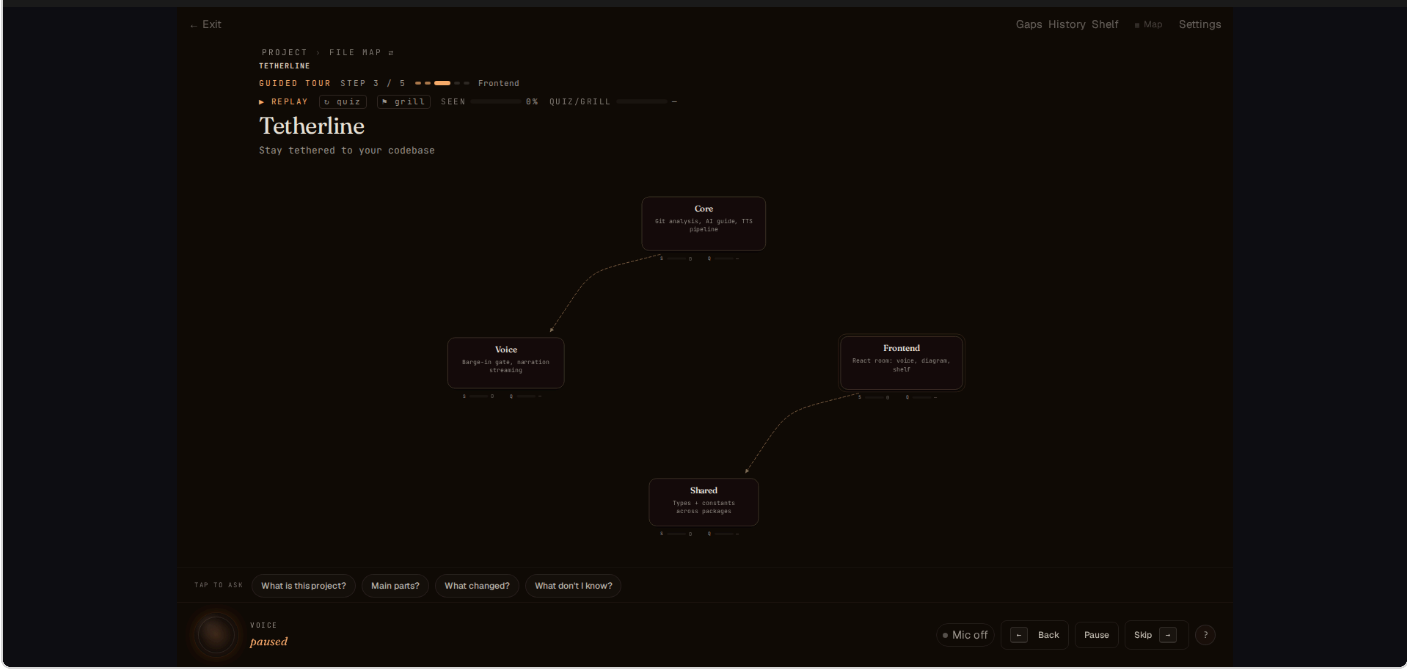

TABLET · 768

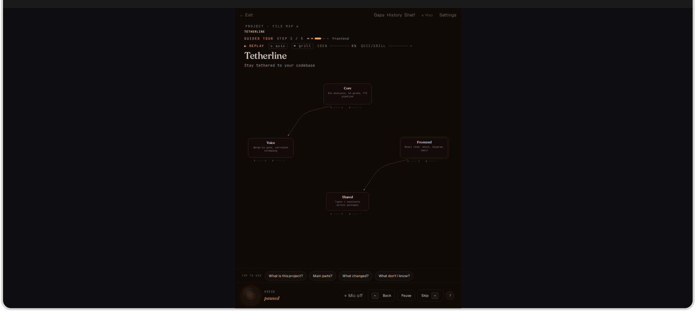

MOBILE · 390

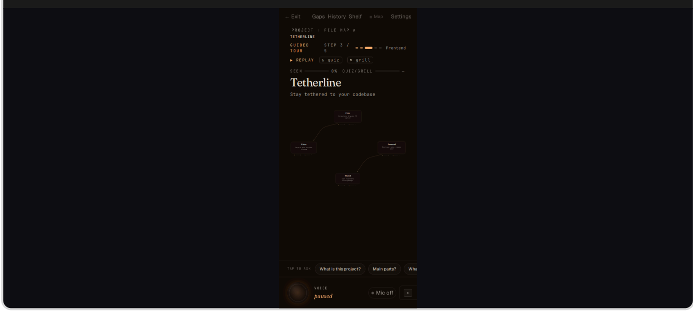

cite as: **p13 · breadcrumb · <viewport>**

You-are-here trail (always-on chrome)

WHO/WHY: not a skill — always-on orientation chrome that appears in the header whenever you are inside a deep_dive pocket, driven by navigation depth. WHAT: the position trail, e.g. CORE › TOKEN REFRESH › 3/8, over the scoped canvas.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

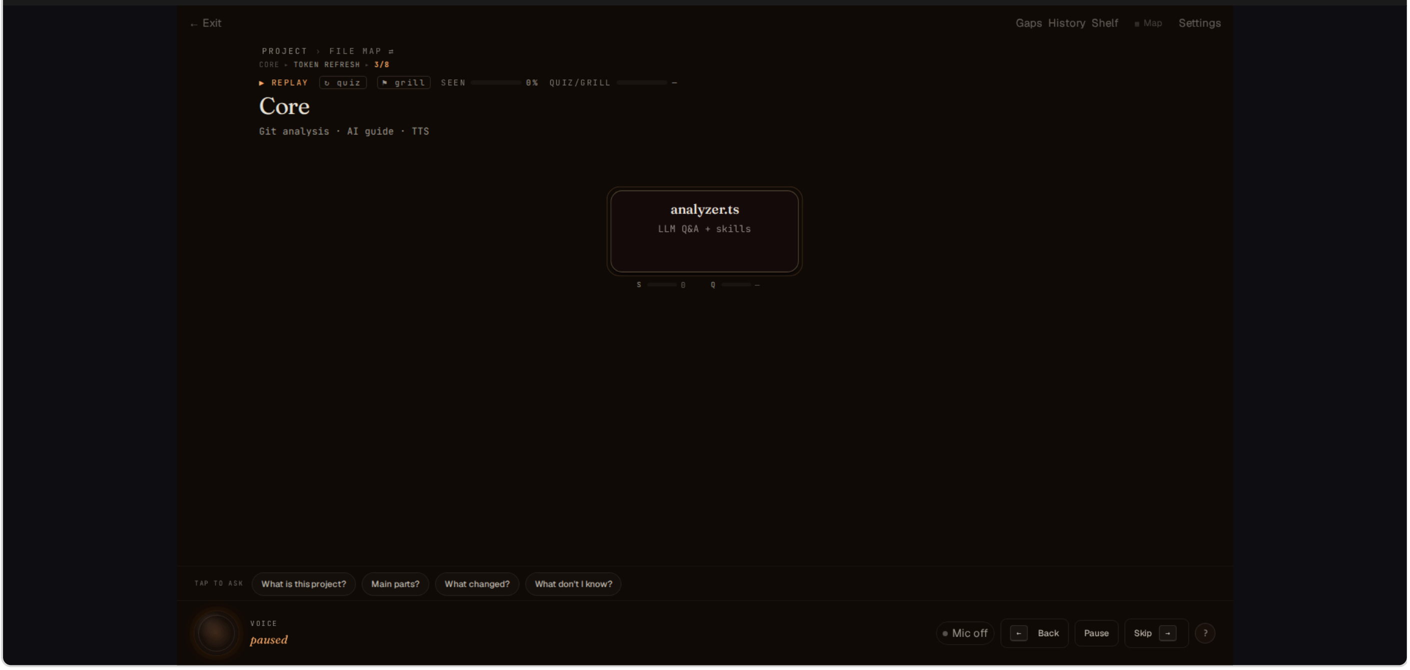

TABLET · 768

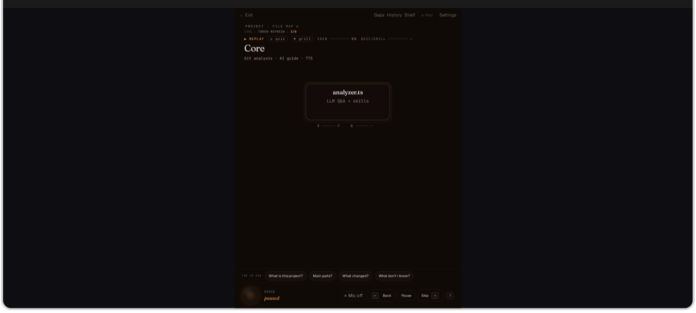

MOBILE · 390

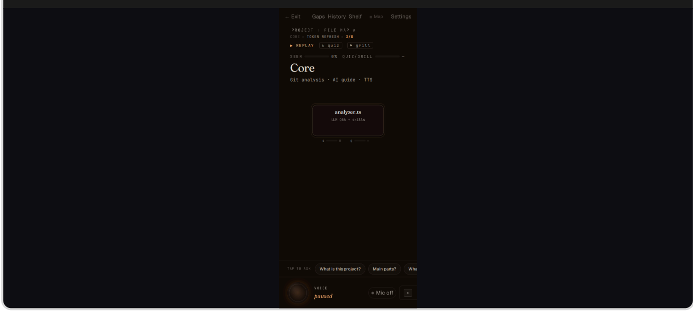

cite as: **p14 · knowledge-layer · <viewport>**

Knowledge — current layer (always-on)

WHO/WHY: not a skill — always-on chrome that renders from live session state on any layer; nobody "calls" it. WHAT: the title block is THIS view's own Seen + best Quiz/Grill (best-for-this-view, not an average); each component shows its rolled-up Seen / Quiz-Grill summary bars with explicit numbers. "Shared S100/Q—" (seen, not proven) is visibly distinct from "Voice S0/Q—" (never seen). Dwell = narration actually finished playing; no verbal-confirmation shortcut.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

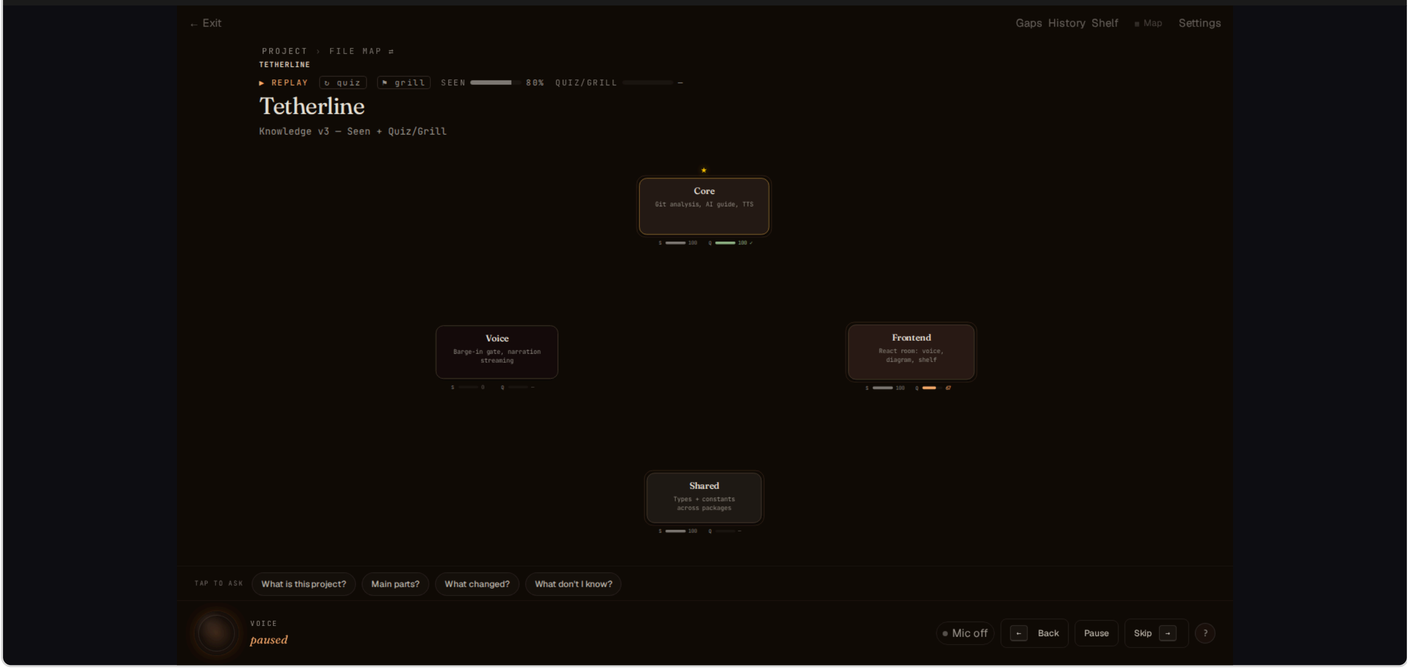

TABLET · 768

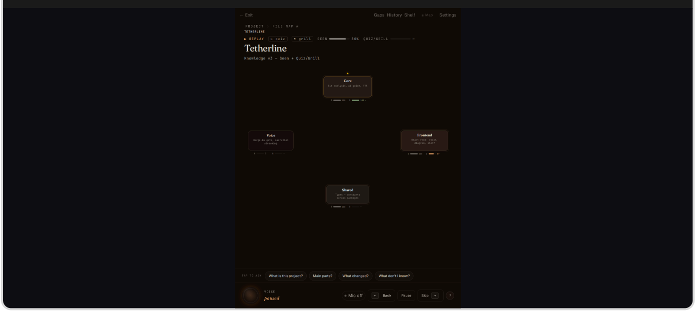

MOBILE · 390

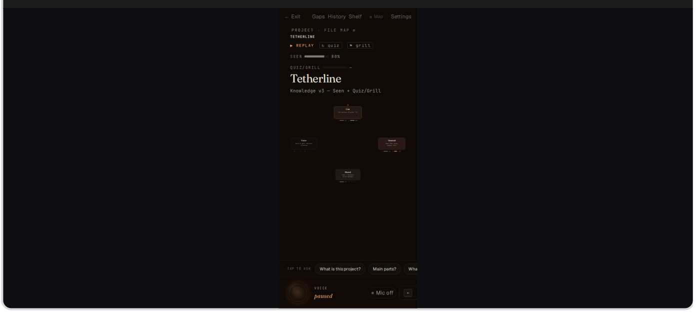

Knowledge — drilled into a module

WHO/WHY: the same always-on knowledge chrome, after you drill into a module. WHAT: the title shows Core's OWN Seen + best quiz (this view only, not the leaf summary); each leaf file shows its own S/Q bars — analyzer.ts S100 Q100 ✓ (grill-passed), chunker.ts S100 Q33, audio-server.py S0 Q— (never seen / never tested).

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

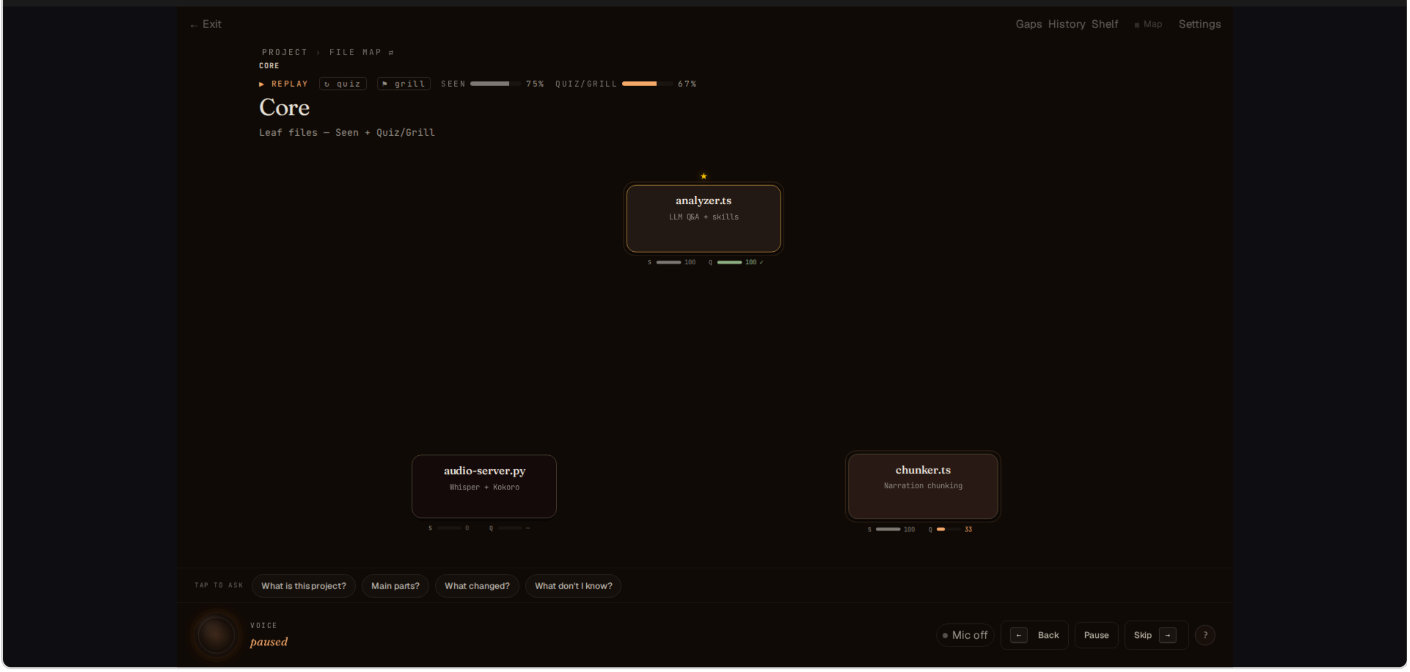

TABLET · 768

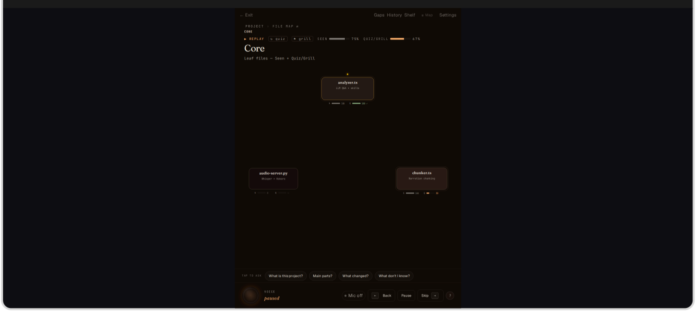

MOBILE · 390

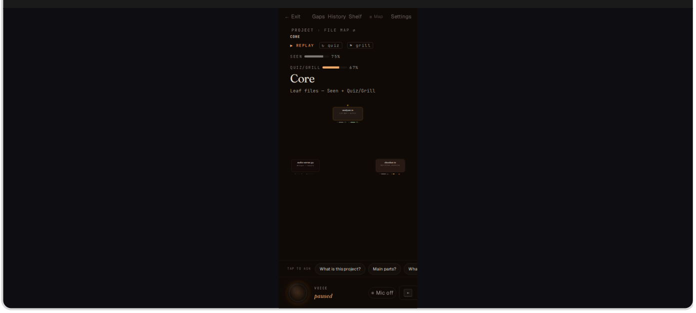

cite as: **p16 · weak-spots-review · <viewport>**

Knowledge — weak-spots review loop

WHO/WHY: YOU open it from the title strip's "weak spots (N)"; it fills automatically as you miss quiz/grill questions. WHAT: every weak/partial question becomes a row with restudy ▶ / resolve ✓ (resolved kept for audit). The actionable pathway to learn more + the supervisor/QA trail.

These shots are rendered from a deterministic UI test fixture (no live backend), so this exact state is reproducible and pixel-checked on every build — that is what the citation id above refers to.

DESKTOP · 1440

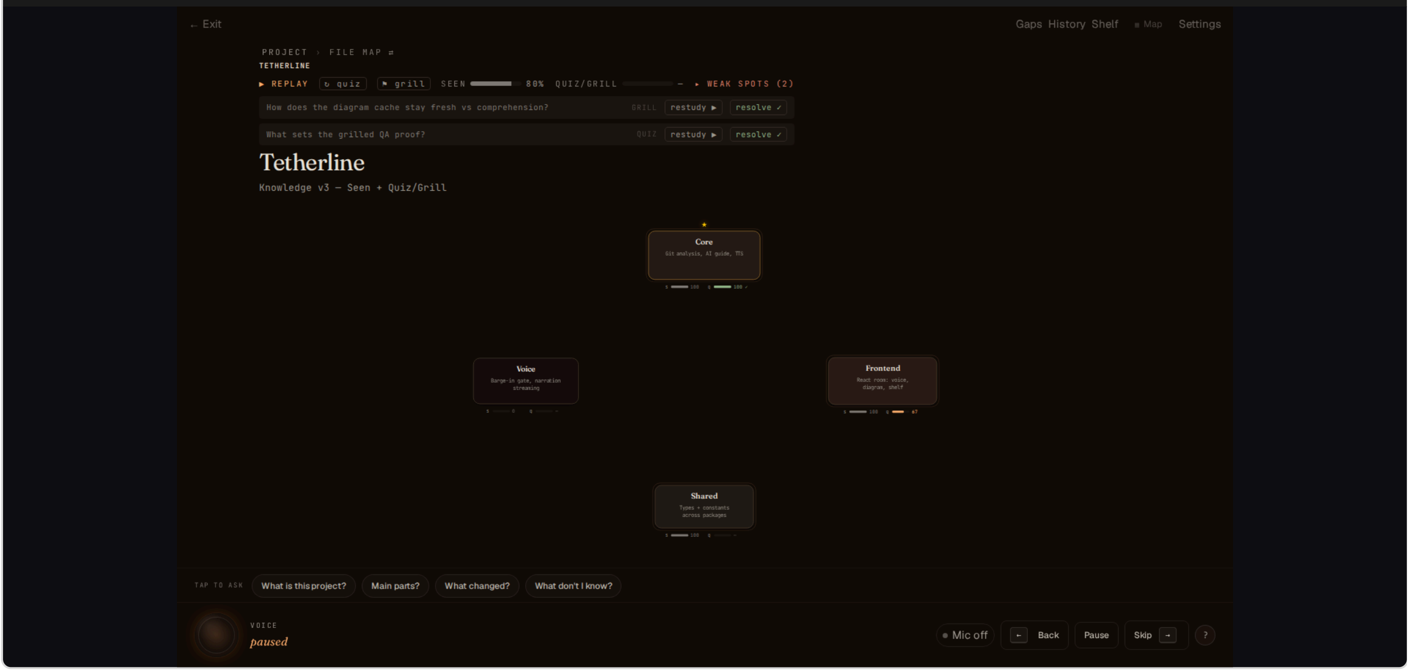

TABLET · 768

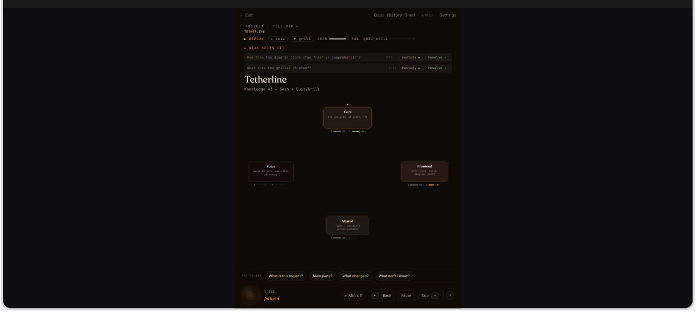

MOBILE · 390